

CPSC 457 - Assignment 3

Due date is posted on D2L.

Individual assignment. Group work is NOT allowed.

Weight: 21% of the final grade.

Both programming questions will ask you to write multithreaded code. You may use pthreads, or C++ threads.

Q1. Programming question – calculating π [10 marks]

Improve the performance of an existing single-threaded `calcp_i` program by converting it to a multi-threaded implementation. Download the starter code, compile it, and run it:

```
$ git clone https://gitlab.com/cpsc457/public/pi-calc.git
$ cd pi-calc
$ make
$ ./calcp_i
Usage: ./calcp_i radius n_threads
       where 0 <= radius <= 100000
       and 1 <= n_threads <= 256
```

The `calcp_i` program estimates the value of π using an algorithm that is very similar to the one described in https://en.wikipedia.org/wiki/Approximations_of_%CF%80#Summing_a_circle's_area. The estimation algorithm is implemented inside the function `count_pi()` in file `calcp_i.cpp`.

The included driver (`main.cpp`) parses the command line arguments, calls `count_pi()` and prints the results. It takes 2 command line arguments: an integer radius and number of threads. For example, to estimate the value of π using radius of 10 and 2 threads, you would run it like this:

```
$ ./calcp_i 10 2
Calculating PI with r=10 and n_threads=2
count: 317
PI:    3.17
```

The function `uint64_t count_pi(int r, int N)` takes two parameters – the radius and number of threads, and returns the number of pixels inside the circle of radius r centered at $(0,0)$ for every pixel (x,y) in square $-r \leq x, y \leq r$. The current implementation is single-threaded, and it ignores the `N` argument. Your job is to re-implement the function so that it uses `N` threads to speed up its execution, such that it runs `N` times faster with `N` threads on hardware where `N` threads can run concurrently. Please note that your assignment will be marked both for correctness and the speedup it achieves.

You need to find a way to parallelize the algorithm **without** using any synchronization mechanisms, such as mutexes, semaphores, atomic types, etc. You are **only allowed** to create and join threads.

I suggest you try to parallelize the outer loop. Give each thread roughly equal number of rows in which to count the pixels. Then sum up the counts from each thread. Your overall algorithm could look like this:

- Create separate memory area for each thread (for input and output), e.g.

```
struct Task { int start_row, end_row, partial_count; ...};
Task tasks[256];
```
- Divide the work evenly between threads, e.g.

```
for(int i = 0 ; i < n_threads ; i ++ ) {
    tasks[i].start_row = ... ;
    tasks[i].end_row = ... ;
}
```
- Create threads and run each thread on the work assigned to it.
- Each thread counts pixels in the rows assigned to it and updates its `partial_count`.
- Join the threads.
- Combine the results of each thread into final result – i.e. return the sum of all `partial_counts`.


Write all code into `calcpi.cpp` and submit this file for grading. Make sure your `calcpi.cpp` works with the included driver program. Your TAs may use a different driver program during marking, so it is important that your code follows the correct API. Make sure your program runs on linuxlab.cpsc.ucalgary.ca.

Assume $0 \leq r \leq 100,000$ and $1 \leq nthreads \leq 256$.

Q2 – Written answer [3 marks]

Time your multi-threaded solution from Q1 with $r = 50000$ using the `time` command on linuxlab.cpsc.ucalgary.ca. Record the real-time for 1, 2, 3, 4, 6, 8, 12, 16, 24 and 32 threads. Also record the timings of the original single-threaded program.

A. Make a table with these timings, and a bar graph, both formatted similar to these:



The numbers in the above table and graph are random. Your timings should look different.

- B. When you run your implementation with N threads, you should see N -times speed up compared to the original single threaded program. Do you observe this in your timings for all values of N ?
- C. Why do you stop seeing the speed up after some value of N ?

Q3. Programming question – detecting primes [30 marks]

Convert a single-threaded program `detectPrimes` to a multi-threaded implementation. Start by downloading the single-threaded code, then compile it and run it:

```
$ git clone https://gitlab.com/cpsc457/public/detectPrimes.git
$ cd detectPrimes/
$ make
$ cat example.txt
    0    3   19 25
4012009 165 1033

$ ./detectPrimes 5 < example.txt
Using 5 threads.
Identified 3 primes:
    3 19 1033
Finished in 0.0000s

$ seq 10000000000000000 10000000000000300 | ./detectPrimes 2
Using 2 threads.
Identified 9 primes:
1000000000000000003 1000000000000000013 1000000000000000019 1000000000000000021
1000000000000000049 1000000000000000081 1000000000000000099 1000000000000000141
1000000000000000181
Finished in 5.6863s
```

The `detectPrimes` program reads integers in range $[2, 2^{63} - 2]$ from standard input, and then prints out the ones that are prime numbers. The first invocation example above detects prime numbers 3, 19 and 1033 in a file `example.txt`. The second invocation uses the program to find all primes in the range $[10^{17}, 10^{17} + 300]$.

`detectPrimes` accepts a single command line argument – a number of threads. This parameter is ignored in the current implementation because it is single-threaded. Your job is to improve the execution time of `detectPrimes` by making it multi-threaded, and your implementation should use the number of threads given on the command line. To do this, you will need to re-implement the function:

```
std::vector<int64_t>
detect_primes(const std::vector<int64_t> & nums, int n_threads);
```

which is defined in `detectPrimes.cpp`. The function takes two parameters: the list of numbers to test, and the number of threads to use. The function is called by the driver (`main.cpp`) after parsing the standard input and command line. Your implementation should use `n_threads` number of threads. Ideally, if the original single-threaded program takes time T to complete a test, then your multi-threaded implementation should finish that same test in T/N time when using N threads. For example, if it takes 10 seconds to complete a test for the original single-threaded program, then it should take your multi-threaded program only 2.5 seconds to complete that same test with 4 threads. To achieve this goal, you will need to design your program so that:

- each thread is doing roughly equal amount of work for all inputs;

- overall your multi-threaded implementation does the same amount of work as the single-threaded version; and
- the synchronization mechanisms you utilize are efficient.

Your TAs will mark your assignment by running the code against multiple different inputs and using different numbers of threads. To get full marks for this assignment, your program needs to:

- output correct results; and
- achieve near optimal speedup for the given number of threads and available cores.

If your code does not achieve optimal speedup on all inputs, you will lose some marks for those tests. Some inputs will include many numbers, some inputs will include just few numbers, some numbers will be large, some small, some will be prime numbers, others will be large composite numbers, etc... For some numbers it will take long time to compute their smallest factor, for others it will take very little time. You need to take all these possibilities into consideration.

Hint 1 – bad solution (do not implement this)

A bad solution would be to parallelize the outer loop of the algorithm, and assign a fixed portion of the numbers to each thread to check. This is a terrible solution because it would not achieve speedups on many inputs, for example where all hard numbers are at the beginning, and all the easy ones at the end. Your program would then likely give all hard numbers to one thread and would end up running just as slowly as the single-threaded version.

Hint 2 – simple solution

A much better, yet still simple solution, would be to parallelize the outer loop, but instead of giving each thread a fixed portion of the input to test, it would decide dynamically how many numbers each thread would process. For example, each thread could process the next number in the list, and if it is a prime, it would add it to the result vector. This would repeat until all numbers have been tested. Note that this solution would achieve optimal speedup for many inputs, but not for all. For example, on input with a single large prime number, it will not achieve any speedup at all. Consequently, if you choose this approach, you will not be able to receive full marks for some tests.

I strongly suggest you start by implementing this simple solution first, and only attempt the more difficult approaches after your simple solution already works.

Hint 3 – very good solution

Even more efficient approach would be to parallelize the inner loop (the loop inside the `is_prime` function). In this approach, all threads would test the same number for primality. If you choose this approach, you need to give each thread a different portion of divisors to check. This will allow you to handle more input cases than the simple solution mentioned earlier. For extra efficiency, and better marks, you also need to consider implementing thread re-use, e.g., by using barriers.

Hint 4 – best solution

This builds on top of the hint 3 above, but it also adds thread cancellation. You need cancellation for cases where one of the threads discovers the number being tested is not a prime, so that it can cancel the work of the other threads. Thread cancellation sounds simple to implement, but it does take considerable effort to get it working correctly.

Write all code into `detectPrimes.cpp` and submit it for grading. Make sure your `detectPrimes.cpp` works with the included `main.cpp` driver. We may use a different driver during marking, so it is important that your code follows the correct API. Make sure your program runs on linuxlab.cpsc.ucalgary.ca.

You may assume that there will be no more than 100,000 numbers in the input, and that all numbers will be in the range $[2, 2^{63} - 2]$.

Please note that the purpose of this question is NOT to find a more efficient factorization algorithm. You must implement the exact same factorization algorithm as given in the skeleton code, except you need to make it multi-threaded.

You may use any of the synchronization mechanisms we covered in lectures, such as semaphores, mutexes, condition variables, spinlocks, atomic variables, and barriers. Don't forget to make sure your code compiles and runs on linuxlab.cpsc.ucalgary.ca.

Q4 – Written question (5 marks)

Time the original single-threaded `detectPrimes.cpp` as well as your multi-threaded version on three files: `medium.txt`, `hard.txt` and `hard2.txt`. For each of these files, you will run your solution 6 times, using 1, 2, 3, 4, 8 and 16 threads. You will record your results in 3 tables, one for each file, formatted like this:

medium.txt			
# threads	Observed timing	Observed speedup compared to original	Expected speedup
original program		1.0	1.0
1			1.0
2			2.0
3			3.0
4			4.0
8			8.0
16			16.0

The 'Observed timing' column will contain the raw timing results of your runs. The 'Observed speedup' column will be calculated as a ratio of your raw timing with respect to the timing of the original single-threaded program. Once you have created the tables, explain the results you obtained. Are the timings what you expected them to be? If not, explain why they differ.

Submission

Submit the following files to D2L for this assignment.

Do not submit a ZIP file. Submit individual files.

<code>calcpi.cpp</code>	solution to Q1
<code>detectPrimes.cpp</code>	solution to Q3
<code>report.pdf</code>	answers to all written questions

Please note – you need to **submit all files every time** you make a submission, as the previous submission will be overwritten.

General information about all assignments:

1. All assignments are due on the date listed on D2L. Late submissions will not be marked.
2. Extensions may be granted only by the course instructor.
3. After you submit your work to D2L, verify your submission by re-downloading it.
4. You can submit many times before the due date. D2L will simply overwrite previous submissions with newer ones. It is better to submit incomplete work for a chance of getting partial marks, than not to submit anything. Please bear in mind that you cannot re-submit a single file if you have already submitted other files. Your new submission would delete the previous files you submitted. So please keep a copy of all files you intend to submit and resubmit all of them every time.
5. Assignments will be marked by your TAs. If you have questions about assignment marking, contact your TA first. If you still have questions after you have talked to your TA, then you can contact your instructor.
6. All programs you submit must run on linuxlab.cpsc.ucalgary.ca. If your TA is unable to run your code on the Linux machines, you will receive 0 marks for the relevant question.
7. Unless specified otherwise, you must submit code that can finish on any valid input under 10s on linuxlab.cpsc.ucalgary.ca, when compiled with `-O2` optimization. Any code that runs longer than this may receive a deduction, and code that runs for too long (about 30s) will receive 0 marks.
8. **Assignments must reflect individual work.** For further information on plagiarism, cheating and other academic misconduct, check the information at this link: <http://www.ucalgary.ca/pubs/calendar/current/k-5.html>.
9. Here are some examples of what you are not allowed to do for individual assignments: you are not allowed to copy code or written answers (in part, or in whole) from anyone else; you are not allowed to collaborate with anyone; you are not allowed to share your solutions with anyone else; you are not allowed to sell or purchase a solution. This list is not exclusive.
10. We will use automated similarity detection software to check for plagiarism. Your submission will be compared to other students (current and previous), as well as to any known online sources. Any cases of detected plagiarism or any other academic misconduct will be investigated and reported.

Appendix 1 – approximate grading scheme for Q3

The test cases that we will use for marking will be designed so that you will get full marks only if you implement the most optimal solution. However, you will receive partial marks even if you one of the less optimal solutions. Here is the rough breakdown of what to expect depending on which solution you implement:

- Parallelization of outer loop with fixed amount of work will yield ~9/28 marks
- Parallelization of outer loop with dynamic work will yield ~15/28 marks
- Parallelization of inner loop without work cancellation and without thread reuse will yield ~19/28 marks
- Parallelization of inner loop with thread reuse (e.g., using barriers) but without cancellation will yield ~24/28 marks.
- Parallelization of inner loop with thread reuse and with cancellation (e.g., using atomic boolean) will give 28/28 marks.

Any tests on which your program produces wrong results will receive 0 marks.